




$\mathcal{S} = \text{local-maxima} \cap \text{thresh}(0.3) \rightarrow 3 \times 3 \text{ dilate} \rightarrow \text{gathered}(b_i, y_i, x_i)$, stock ESOD

Blue-green = SparseHead: inference-only, re-uses dense-trained weights, no retraining

Skipped positions are never computed (not computed-then-discarded)

Composes with any selector arm; offsets a larger retained area (Occupy) at inference

Not enabled for the numbers in Table 1/2 of the paper – flagged as the next lever to test